

LLMs for Multi-Agent System Control Project Proposal

Noah Cortes
noah.cortes@gatech.edu

Richard Koulen
rkoulen3@gatech.edu

Nathaniel Wert
nwert3@gatech.edu

1 Motivation

Multi-agent systems (MAS) have vast potential for various applications of robotics. From search and rescue [1], to warehouse logistics [2] MAS can be used to improve efficiency and effectiveness in various domains. However, one challenge with MAS is the significant overhead in deployment and developing such systems.

Because deployment of MAS is often a complex and expensive task, most of the body of research in MAS is completed in simulation. Specifically, a large body of research exists in real-time strategy games (specifically StarCraft) [3]. Within this body of research, the game defines a set of constraints and actions that agents can take while providing reinforcement learning algorithms significant feedback on the performance of the agents [4]. Because of the success of training agents in games, there has become a growing interest in applying these techniques to more realistic and increasingly complex scenarios [5]. However, training agents in games is still complex and expensive and many real-world deployments require a human in the loop to ensure that the agents are performing as expected. This is especially true in scenarios where the agents are operating in a dynamic environment with changing constraints and objectives.

Branching out of the need for a human in the loop, there has been a growing interest in using large language models (LLMs) to control agents in MAS. In most research, this LLM control has taken the form of high level task allocation, mid level motion planning, and low level action generation [6]. For this project, we will be focusing on the high level task allocation problem. This problem has increasingly been shown to be accomplishable by LLMs, with LLMs being able to generate task allocation plans for MAS in a variety of standard MAS and swarm problem formulations (such as flocking and formation control) [7], [8], [9] as well as robot soccer [10]. Although many different approaches have been attempted for LLM control of MAS, a particularly interesting approach involves having LLMs generate plans using a sort of intermediate representation. In [11], the authors have the LLM generate a behavior tree for the agents to follow. In [12], the authors use Linear Temporal Logic (LTL) as an intermediate representation for the LLM to generate plans.

While LLMs have been shown to be able to generate plans for MAS, there has been little work done in evaluating these systems with a human in the loop and in situations with many robots. First, much of the previous work has focused on either pre-defined and common MAS problems (such as flocking and formation control) [9] or on a small number of robots without a human in the loop [7], [8], [13]. Second, there has been little work done in evaluating the user-side metrics of these systems. In [12], the authors conducted a 4 participant user study to evaluate the usability of their system. However, this study was limited to the performance of the system and did not evaluate the user experience of the participants.

Therefore, in this project, we will be evaluating and comparing the performance and user-experience of a traditional real-time strategy game interface and a LLM interface for controlling a MAS in a simulated game environment. Specifically, we are interested in investigating whether utilizing an LLM interface for controlling the MAS leads to decreased task load and what impact the LLM interface has on performance. Additionally, much of the previous research on using LLMs to control MAS has focused on contrived examples so we are additionally interested in investigating how humans interact and perceive the LLM interface when controlling a MAS in a more complex scenario.

2 Proposed Methodology

The main goal of our project is to create a real-time strategy game interface and an LLM interface for controlling a MAS in a game environment. We will then perform a **within-subjects** user study to evaluate the performance and user experience of the two interfaces (**counterbalanced** against the interface selection). Specifically, we will be evaluating the performance of the MAS in terms of the score achieved in the game and the user experience of the participants in terms of task load and usability. To measure task load, we will be using the **NASA Task Load Index (NASA-TLX)** [14]. To measure usability, we will be using the **System Usability Scale (SUS)** [15]. Additionally, we will be collecting qualitative feedback from the participants to gain insight into their experience with the two interfaces.

The first major part of our project is creating a game environment for the MAS to operate in. We will be using the Robotarium [16] as our MAS platform and will be creating a game environment for the MAS to operate within. Our main goals for the game environment are to create a game that is challenging enough to require coordination of a large number of agents and that includes a both positive and negative rewards. Specifically, much of the previous work has focused on contrived examples that do not require conforming to a set of constraints, whereas these are incredibly important aspects of real-world MAS problems. Therefore, our game environment will include a set of negative feedback constraints that must be communicated to the MAS by the human operator. Additionally, we will be creating a scoring system for the game that will provide positive feedback to the MAS for completing tasks and negative feedback for failing to complete tasks or violating constraints.

With the Robotarium providing a software simulator, hardware, and runtime environment for the MAS, we will be able to focus on creating the game environment and interfaces for controlling the MAS. The team has not yet decided on the specific LLM for use in the project, but selecting the LLM will be one of the first tasks for the project. Additionally, we will be creating a traditional real-time strategy game interface for controlling the MAS. A technology stack has also not yet been selected for the real-time strategy game interface, but the design is meant to adhere to the human-computer interaction principles for real-time strategy games as presented in [17] and [18].

3 Proposed Evaluation

The project is considered successful if the following criteria are met:

- A game environment is created utilizing the Robotarium platform
- A real-time strategy game interface is created for controlling the MAS
- An LLM interface is created for controlling the MAS
- A user study is conducted to evaluate the performance and user experience of the two interfaces

4 Project Timeline

4.1 Project Check-In Day 1

By project check-in day 1, the following milestones should ideally be completed:

- The game environment should be established and built within the Robotarium platform.
- The real-time strategy game interface should be designed and partially implemented
- The LLM interface should be designed and partially implemented

4.2 Project Mid-Point "Poster" Session

By the project mid-point poster session, the following milestones should ideally be completed (in addition to the milestones above):

- The real-time strategy game interface should be fully implemented and functional
- The LLM interface should be fully implemented and functional
- The user study should be designed and ready to be conducted
- The mid-point poster should be completed and presented

4.3 Project Check-In Day 2

By project check-in day 2, the following milestones should ideally be completed (in addition to the milestones above):

- The user study should be conducted and completed

4.4 Final Project Presentations

By the Final project presentation, the following milestones should ideally be completed (in addition to all other milestones above):

- The user study should be analyzed and the results should be presented
- The final project poster should be completed and presented

5 Proposed Grading Rubric

Proposed Grading Rubric	
Finding / System Component / Capability / Validation Test	Points
Game Design and Implementation	30
Real-Time Strategy Interface	30
LLM Interface	30
User Study Setup	30
User Study Data Analysis and Interpretation	30

References

- [1] D. S. Drew. Multi-Agent Systems for Search and Rescue Applications. *Current Robotics Reports*, 2(2):189–200, June 2021. ISSN 2662-4087. doi:10.1007/s43154-021-00048-3. URL <https://doi.org/10.1007/s43154-021-00048-3>.
- [2] P. O. Dusadeerungsikul, X. He, M. Sreeram, and S. Y. Nof. Multi-agent system optimisation in factories of the future: cyber collaborative warehouse study. *International Journal of Production Research*, 60(20):6072–6086, Oct. 2022. ISSN 0020-7543. doi:10.1080/00207543.2021.1979680. URL <https://doi.org/10.1080/00207543.2021.1979680>.
- [3] Jensen Andreas Schmidt, Kaysørdam Christian, and Villadsen Jørgen. Interfacing Agents to Real-Time Strategy Games. In *Frontiers in Artificial Intelligence and Applications*. IOS Press, 2015. doi:10.3233/978-1-61499-589-0-68. URL <https://www.medra.org/servlet/aliasResolver?alias=iospressISBN&isbn=978-1-61499-588-3&spage=68&doi=10.3233/978-1-61499-589-0-68>.
- [4] W. J. Zhou, B. Subagdja, A.-H. Tan, and D. W.-S. Ong. Hierarchical control of multi-agent reinforcement learning team in real-time strategy (RTS) games. *Expert Systems with Applications*, 186:115707, Dec. 2021. ISSN 09574174. doi:10.1016/j.eswa.2021.115707. URL <https://linkinghub.elsevier.com/retrieve/pii/S0957417421010897>.
- [5] V. G. Goecks, N. Waytowich, D. E. Asher, S. Jun Park, M. Mittrick, J. Richardson, M. Vindiola, A. Logie, M. Dennison, T. Trout, P. Narayanan, and A. Kott. On games and simulators as a platform for development of artificial intelligence for command and control. *The Journal of Defense Modeling and Simulation*, 20(4):495–508, Oct. 2023. ISSN 1548-5129. doi:10.1177/15485129221083278. URL <https://doi.org/10.1177/15485129221083278>.
- [6] P. Li, Z. An, S. Abrar, and L. Zhou. Large language models for multi-robot systems: a survey. *Autonomous Robots*, 50(3):30, June 2026. ISSN 1573-7527. doi:10.1007/s10514-026-10257-4. URL <https://doi.org/10.1007/s10514-026-10257-4>.
- [7] Y. Wang, N. Kemsaram, and S. Subramanian. Decentralized LLM-Driven Coordination of Acoustic Robots for Contactless Object Manipulation, May 2026. URL <http://arxiv.org/abs/2605.29378>. arXiv:2605.29378 [cs.RO].
- [8] E. M. Eumi, H. Abbass, and N. Marcus. SwarmChat: An LLM-Based, Context-Aware Multimodal Interaction System for Robotic Swarms. In Y. Tan and Y. Shi, editors, *Advances in Swarm Intelligence*, pages 181–192, Singapore, 2026. Springer Nature. ISBN 9789819509829. doi:10.1007/978-981-95-0982-9_15.
- [9] W. Ji, H. Chen, M. Chen, G. Zhu, L. Xu, R. Groß, R. Zhou, M. Cao, and S. Zhao. GenSwarm: Scalable Multi-Robot Code-Policy Generation and Deployment via Language Models. *npj Robotics*, 4(1):5, Jan. 2026. ISSN 2731-4278. doi:10.1038/s44182-025-00065-w. URL <https://www.nature.com/articles/s44182-025-00065-w>.
- [10] M. Brienza, E. Musumeci, V. Suriani, D. Affinita, A. Pennisi, D. Nardi, and D. D. Bloisi. LLCoach: Generating Robot Soccer Plans Using Multi-role Large Language Models. In E. Barros, J. P. Hanna, H. Okada, and E. Torta, editors, *RoboCup 2024: Robot World Cup XXVII*, pages 176–188, Cham, 2025. Springer Nature Switzerland. ISBN 978-3-031-85859-8. doi:10.1007/978-3-031-85859-8_15.
- [11] M. Majid and A. Y. Majid. CommandSwarm: Safety-Aware Natural Language-to-Behavior-Tree Generation for Robotic Swarms, May 2026. URL <http://arxiv.org/abs/2605.07764>. arXiv:2605.07764 [cs.RO].
- [12] S. Xu, X. Luo, Y. Huang, L. Leng, R. Liu, and C. Liu. NI2Hlt2Plan: Scaling Up Natural Language Understanding for Multi-Robots Through Hierarchical Temporal Logic Task

- Specifications. *IEEE Robotics and Automation Letters*, 10(10):10482–10489, Oct. 2025. ISSN 2377-3766. doi:10.1109/LRA.2025.3598648. URL <https://ieeexplore.ieee.org/document/11123780>.
- [13] S. Borate, B. R. B. V. Pardeshi, and M. Vadali. CoMuRoS - An LLM-based generalizable hierarchical task planning and execution framework for heterogeneous robot teams with event-driven re-planning. *Frontiers in Robotics and AI*, 13, Aug. 2026. ISSN 2296-9144. doi:10.3389/frobt.2026.1843313. URL <https://www.frontiersin.org/journals/robotics-and-ai/articles/10.3389/frobt.2026.1843313/full>.
 - [14] S. G. Hart and L. E. Staveland. Development of NASA-TLX (Task Load Index): Results of empirical and theoretical research. In *Human mental workload*, Advances in psychology, 52, pages 139–183. North-Holland, Oxford, England, 1988. ISBN 978-0-444-70388-0. doi:10.1016/S0166-4115(08)62386-9.
 - [15] J. Brooke. SUS-A quick and dirty usability scale. *Usability evaluation in industry*, 189(194), 1996.
 - [16] D. Pickem, P. Glotfelter, L. Wang, M. Mote, A. Ames, E. Feron, and M. Egerstedt. The Robotarium: A remotely accessible swarm robotics research testbed. In *2017 IEEE International Conference on Robotics and Automation (ICRA)*, pages 1699–1706, May 2017. doi:10.1109/ICRA.2017.7989200. URL <https://ieeexplore.ieee.org/document/7989200>.
 - [17] H. Wenjun and B. Xiudong. HCI in real-time strategy games: a study of principals and guidelines for designing 3D user interface. In *2006 7th International Conference on Computer-Aided Industrial Design and Conceptual Design*, pages 1–6, Nov. 2006. doi:10.1109/CAIDCD.2006.329386. URL <https://ieeexplore.ieee.org/document/4127155>.
 - [18] M. S. B. Riaz and A. Javed. User Interface Designing Principles for Real-time Games Strategies. *Journal of Computing & Biomedical Informatics*, 2(01):128–139, Mar. 2021. ISSN 2710-1614. doi:10.56979/201/2021/29. URL <https://jcbi.org/index.php/Main/article/view/29>.